

Hadoop MapReduce

Claudio Scheer¹ Gabriell Araujo¹

¹Master's Degree in Computer Science
Pontifical Catholic University of Rio Grande do Sul - PUCRS

High Performance for Big Data Applications

Table of Contents

1 Hadoop

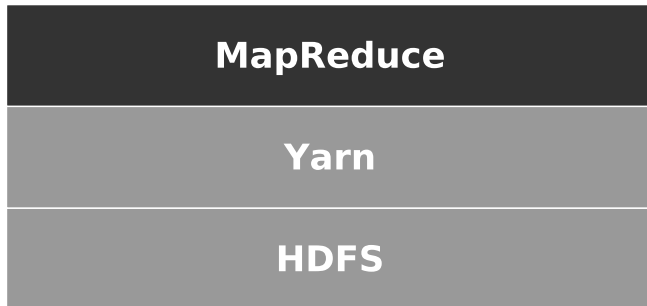
2 Examples

Table of Contents

1 Hadoop

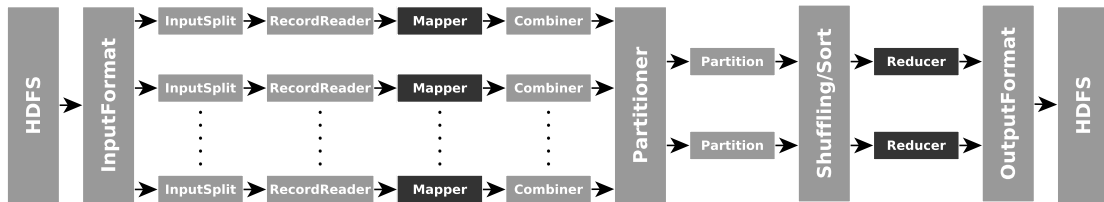
2 Examples

What we mean by Hadoop



<https://data-flair.training/blogs/hadoop-ecosystem-components>

MapReduce execution flow



<https://data-flair.training/blogs/hadoop-ecosystem-components>

Custom data types

- LongWritable = long;
- IntWritable = int;
- Text = String;
- Other data types ([link](#));

```
public class IntWritable implements WritableComparable<IntWritable> {
    private int value;

    public IntWritable(int value) { set(value); }

    public void set(int value) { this.value = value; }

    public int get() { return value; }

    @Override
    public void readFields(DataInput in) throws IOException {
        value = in.readInt();
    }

    @Override
    public void write(DataOutput out) throws IOException {
        out.writeInt(value);
    }

    @Override
    public int compareTo(IntWritable o) {
        int thisValue = this.value;
        int thatValue = o.value;
        return (thisValue < thatValue ? -1 : (thisValue == thatValue ? 0 : 1));
    }
}
```

► Source

- TextInputFormat: <LongWritable, Text>
- KeyValueTextInputFormat: <Text, Text>
 - Key splitted by \t;
- NLineInputFormat: <LongWritable, Text>
 - `config.setInt(NLineInputFormat.LINES_PER_MAP, 256);`
- Customs InputFormat must implement `getSplits` and `getRecordReader`;
 - <https://hadoop.apache.org/docs/stable/api/org/apache/hadoop/mapred/InputFormat.html>

- Defines the parallelism level;
- Usually splitted by the size of the block;
 - $10\text{TB}/128\text{MB} = 82000$
- It is a logical division of the input data;
- Ways to change it:
 - `config.set(MRJobConfig.NUM_MAPS, 2);`
 - `mapreduce.job.maps` on `mapsite-site.xml`;

<https://hadoop.apache.org/docs/stable/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html>

- Split InputSplit into <key, value> pairs;
 - Custom implementations can read values from outside the InputSplit;
- Max record length:
 - `config.setInt(LineRecordReader.MAX_LINE_LENGTH, Integer.MAX_VALUE);`
- Vanilla example:

```
run(Context context) {  
    while(context.nextKeyValue())  
    {  
        map(context.setCurrentKey(), context.getCurrentValue(), context)  
    }  
}
```

- Maps $\langle \text{key1}, \text{value1} \rangle$ to $\langle \text{key2}, \text{value2} \rangle$;
- Vanilla example:

```
public class SimpleMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    private IntWritable one = new IntWritable(1);  
    private Text word = new Text();  
  
    @Override  
    public void map(LongWritable key, Text value, Context context) {  
        StringTokenizer itr = new StringTokenizer(value.toString());  
        while (itr.hasMoreElements()) {  
            word.set(itr.nextToken());  
            context.write(word, one);  
        }  
    }  
}
```

Combiner

- Reduces data transfer between mapper and reducer;
- Reduces the amount of data to be processed in the reducer;



Partitioner

- Redirects the Combiner output to a specific reducer;
- Has the same number as the number of reducers;
- Used only when there is more than one reducer;
- Vanilla example:

```
public class StupidPartitioner extends Partitioner<Text, IntWritable> {  
    public int getPartition(Text key, IntWritable value, int numPartitions) {  
        if (value.get() < 35) {  
            return 0;  
        } else {  
            return 1;  
        }  
    }  
}
```

<https://hadoop.apache.org/docs/stable/api/org/apache/hadoop/mapreduce/Partitioner.html>

- Collects output from the mappers to the reducers using HTTP requests;
- Sorts the collected <key, value> pairs based on the key;

<https://hadoop.apache.org/docs/stable/hadoop-mapreduce-client/hadoop-mapreduce-client-core/PluggableShuffleAndPluggableSort.html>

<https://hadoop.apache.org/docs/stable/hadoop-mapreduce-client/hadoop-mapreduce-client-core/EncryptedShuffle.html>

- Maps $\langle \text{key2}, \text{list}(\text{value2}) \rangle$ to $\langle \text{key3}, \text{value3} \rangle$;
- Vanilla example:

```
public static class SimpleReducer extends Reducer<Text, IntWritable, Text, IntWritable> {  
    private IntWritable result = new IntWritable();  
  
    @Override  
    public void reduce(Text key, Iterable<IntWritable> values, Context context) {  
        int sum = 0;  
        for (IntWritable val : values) {  
            sum += val.get();  
        }  
        result.set(sum);  
        context.write(key, result);  
    }  
}
```

- Specifies how reducer output will be written;
- TextInputFormat: `<LongWritable, Text>`
- KeyValueTextInputFormat: `<Text, Text>`
 - Key splitted by `\t`;
- NLineInputFormat: `<LongWritable, Text>`
 - `config.set(TextOutputFormat.SEPARATOR, “,”);`
- Custom InputFormat must implement `getSplits` and `getRecordReader`;
 - <https://hadoop.apache.org/docs/stable/api/org/apache/hadoop/mapred/InputFormat.html>

Table of Contents

1 Hadoop

2 Examples

- WordCount; - CountProductsSold; -

Sample frame title

In this slide, some important text will be highlighted because it's important. Please, don't abuse it.

Remark

Sample text

Important theorem

Sample text in red box

Examples

Sample text in green box. The title of the block is "Examples".

This is a text in first column.

$$E = mc^2$$

- First item
- Second item

This text will be in the second column and on a second thought this is a nice looking layout in some cases.